

1 Overview

The goal of this project is to train an agent with a double-jointed arm to follow a sphere. The agent receives a reward of +0.1 in every step where the hand of the agent is inside the sphere. This example can be easily transferred to the real world, where it is a common challenge to program a robot arm to reach a certain point in space.

For training the agent I used a DDPG (Deep Deterministic Policy Gradient) algorithm. DDPG was chosen over Q learning because it performs way better in a problem with continuous action spaces. Nevertheless, this Project uses multiple ideas from Q learning in addition to the DDPG algorithm:

Replay memory:

The agent utilizes the experience replay technique during training. It stores each experience in a buffer. In the learning step the agent will query a random sample from the buffer and learns from those randomly sampled experiences.

Target networks:

When updating the weights of the neural network of the agent we not only update the value for $Q(s, a)$ but also of other state action pairs $Q(s', a')$. This can lead to a very unstable training of the agent. In order to stabilize the training we use two neural networks:

The **local network** is trained at every step to get the best prediction for the Q-value.

The agent uses the **target network** to choose actions. In this project the target network is updated by Soft Update from the local network.

Additionally the DDPG algorithm introduces the actor/critic technique. The critic receives a state, and an action as input and estimates the expected reward for this state, action tuple. The actor receives a state and outputs an action vector with each value corresponding to a continuous action.

The state space consists of 33 continuous states from which the agent estimates 4 continuous action values. Both networks (actor & critic) consist of three fully connected layers with a node size of 128 for the first and second layer. However, the Actor network uses a tanh activation function for the last layer to get a value between -1 and 1.

2 Results

This model took a long time to train...

3 Ideas for future improvements

Parallelize the training to multiple agents at the same time. This could for example be implemented with the advantage actor critic (A2C) Algorithm where multiple agents share the same network.